

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from prepare import wrangle_data, split_data
import explore
import model

import warnings
warnings.filterwarnings('ignore')

target = 'newconstructionyn'
```

**Note** all code for this project can be found in the Github repository at [github.com/bensmith07/new\\_construction](https://github.com/bensmith07/new_construction)

## Project Description

This project uses a labeled set of MLS listings from San Antonio, TX and attempts to build a model that will identify which of those listings are new construction and which are not. For a company which aims to provide the most accurate Automated Valuation Model (AVM) on the market, the utility of such a prediction is obvious. The market for new construction is likely to be different than that of existing homes, and we should expect a requisite difference in value between a newly constructed home and a similar home which has been previously owned.

### Assumptions

For the purposes of this analysis, I have made the assumption that the ultimate goal for the model will be to identify new construction for a new listing (at the time the property is listed) rather than to retroactively identify new construction amidst a set of historical listings.

I have also made the assumption that the model should be generalizable to other metropolitan areas.

Based on these assumptions, I removed data which:

- A) would not be available or would be irrelevant at the time a property is first listed on the MLS
- B) identified specific geographic localities

I also assume that the ultimate goal is to identify as many new construction listings as possible, and to avoid letting a new construction listing go by unnoticed. This assumption affected the metric by which I chose to evaluate predictive models.

### Goals

- 1) Create a model to identify new construction vs. existing homes
- 2) Summarize the aggregate difference in price between new construction vs. existing homes

## Acquisition

This dataset came as a prepackaged CSV file containing the features shown below

In [2]:

```
df = pd.read_csv('data.csv')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4780 entries, 0 to 4779
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   address_id                           4780 non-null   int64
1   city                                 4778 non-null   object
2   postal_code                          4778 non-null   float64
3   census_tract                        4778 non-null   float64
4   standardstatus                      4780 non-null   object
5   mlsstatus                           4780 non-null   object
6   contractstatuschangedate            4780 non-null   object
7   purchasecontractdate                3922 non-null   object
8   closedate                           3508 non-null   object
9   listingcontractdate                 4780 non-null   object
10  daysonmarket                        4779 non-null   float64
11  closeprice                          3508 non-null   float64
12  listprice                           4780 non-null   int64
13  originallistprice                   4765 non-null   float64
14  propertytype                        4780 non-null   object
15  propertysubtype                     4779 non-null   object
16  lotsizearea                         4675 non-null   float64
17  lotsizeunits                        4612 non-null   object
18  lotfeatures                          1243 non-null   object
19  totalactualrent                     1182 non-null   float64
20  bedroomtotal                        4764 non-null   float64
21  bathroomtotalinteger                4764 non-null   float64
22  bathroomfull                       4764 non-null   float64
23  bathroomshalf                      4764 non-null   float64
24  livingarea                          4733 non-null   float64
25  livingareaunits                     4733 non-null   object
26  garageyn                            4727 non-null   float64
27  parkingfeatures                     4764 non-null   object
28  stories                             3588 non-null   float64
29  yearbuilt                           4772 non-null   float64
30  newconstructionyn                  4328 non-null   float64
31  heating                             4770 non-null   object
32  cooling                              4770 non-null   object
33  architecturalstyle                  4760 non-null   object
dtypes: float64(16), int64(2), object(16)
memory usage: 1.2+ MB
```

## Preparation

In order to clean and prepare this data, I performed the following steps:

### Removing Irrelevant or Un-Useful Data

Including:

- any column where there is only one unique value
- location specific data
- data unavailable/irrelevant at the time of a new listing

### Handling Missing Values

- removing columns with an abundance of null values ( `lotfeatures` , `totalactualrent` )
- imputing missing values for the following columns:
  - `garageyn` (imputed based on information found in `parkingfeatures` column)
  - `stories` (imputed based on information found in `architecturalstyle` column)
  - `lotsizearea` (imputed with the median)

### Adjusting Data Types

- `address_id`: `int` >> `string`
- `listingcontractdate`: `string` >> `Pandas Timestamp`

### Data Cleaning

- `propertysubtype` : combining category labels which likely have the same meaning ('Single Family Detached' and 'Single Family Residence Detached')
- `lotsizearea` : correcting negative lot size values (which were likely listed in error) by taking the absolute value

### Engineering New Features:

including:

(for the sake of time, I have excluded explicit definitions of these features)

- `listing_month`
- `listing_dayofmonth`
- `listing_dayofweek`
- `listed_on_weekend`
- `years_since_build`
- `built_last_two_years`
- `parkingfeatures_attached`
- `parkingfeatures_detached`
- `parkingfeatures_oversized`

- parkingfeatures\_converted
- parkingfeatures\_sideentry
- parkingfeatures\_reareentry
- parkingfeatures\_tandem
- parkingfeatures\_golfcart
- heating\_central
- heating\_naturalgas
- heating\_electric
- heating\_heatpump
- heating\_2units
- heating\_1unit
- heating\_zoned
- heating\_other
- heating\_floorfurnace
- heating\_solar
- heating\_propaneowned
- heating\_none
- heating\_windowunit
- cooling\_central
- cooling\_winddownwall
- cooling\_heatpump
- cooling\_zoned
- central\_cooling\_units
- winddownwall\_cooling\_units
- archstyle\_traditional
- archstyle\_splitlevel
- archstyle\_ranch
- archstyle\_texashillcountry
- archstyle\_craftsman
- archstyle\_other
- archstyle\_colonial
- archstyle\_spanish
- archstyle\_manufacturedhome-singlewide
- archstyle\_a-frame
- lotsizearea\_listed\_negative
- lotsizearea\_small
- originallistprice\_persqft
- originallistprice\_scaled
- originallistprice\_scaled\_persqft
- previously\_listed

### **For-Sale vs. For-Rent Listings**

Seeing that the `originallistprice` of For-Rent represents the monthly rent price, whereas in For-Sale listings it represents the purchase price, and that these two things are on entirely different scales, I created Min-Max scaled versions of the `originallistprice` and `originallistprice_persqft` column.

However, the distributions of these scaled columns were still quite different from one another, with the rental properties having `originallistprice_persqft_scaled` values that were much more heavily concentrated around the mean.

Recognizing that this could have a negative affect on any modeling and analysis, I created two separate dataframes, one for rental listings and one for sale listings, which could be used to analyze or model these categories independent of each other, while also maintaining the original whole dataset in it's own dataframe variable.

(Note: Analysis and modeling with the separate dataframes is a step I would take if given more time for this project, but this iteration of the project does not utilize these separate

In [3]:

```
# this function performs the preparation steps described above
df, sale_df, rent_df = wrangle_data()
# displaying info about the transformed dataframe
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 4289 entries, 0 to 4288
```

```
Data columns (total 66 columns):
```

#	Column	Non-Null Count	Dtype
0	address_id	4289 non-null	object
1	listingcontractdate	4289 non-null	datetime64[ns]
2	originallistprice	4289 non-null	float64
3	propertytype	4289 non-null	object
4	propertysubtype	4289 non-null	object
5	lotsizearea	4289 non-null	float64
6	bedroomstotal	4289 non-null	float64
7	bathroomstotalinteger	4289 non-null	float64
8	bathroomsfull	4289 non-null	float64
9	bathroomshalf	4289 non-null	float64
10	livingarea	4289 non-null	float64
11	garageyn	4289 non-null	bool
12	stories	4289 non-null	float64
13	yearbuilt	4289 non-null	float64
14	newconstructionyn	4289 non-null	bool
15	listing_month	4289 non-null	int64
16	listing_dayofmonth	4289 non-null	int64
17	listing_dayofweek	4289 non-null	int64
18	listed_on_weekend	4289 non-null	bool
19	previously_listed	4289 non-null	bool
20	years_since_build	4289 non-null	float64
21	built_last_two_years	4289 non-null	bool
22	garage_size	4289 non-null	int64
23	parkingfeatures_attached	4289 non-null	bool
24	parkingfeatures_detached	4289 non-null	bool
25	parkingfeatures_oversized	4289 non-null	bool
26	parkingfeatures_converted	4289 non-null	bool

```

27 parkingfeatures_sideentry          4289 non-null    bool
28 parkingfeatures_reareentry         4289 non-null    bool
29 parkingfeatures_tandem              4289 non-null    bool
30 parkingfeatures_golfcart            4289 non-null    bool
31 heating_central                     4289 non-null    bool
32 heating_naturalgas                  4289 non-null    bool
33 heating_electric                    4289 non-null    bool
34 heating_2units                      4289 non-null    bool
35 heating_heatpump                    4289 non-null    bool
36 heating_1unit                       4289 non-null    bool
37 heating_zoned                       4289 non-null    bool
38 heating_other                       4289 non-null    bool
39 heating_floorfurnace                4289 non-null    bool
40 heating_solar                       4289 non-null    bool
41 heating_propaneowned                4289 non-null    bool
42 heating_none                        4289 non-null    bool
43 heating_windowunit                 4289 non-null    bool
44 cooling_central                     4289 non-null    bool
45 cooling_windowwall                  4289 non-null    bool
46 cooling_heatpump                    4289 non-null    bool
47 cooling_zoned                       4289 non-null    bool
48 central_cooling_units               4289 non-null    int64
49 windowwall_cooling_units            4289 non-null    int64
50 archstyle_traditional                4289 non-null    bool
51 archstyle_contemporary              4289 non-null    bool
52 archstyle_splitlevel                4289 non-null    bool
53 archstyle_ranch                     4289 non-null    bool
54 archstyle_texashillcountry           4289 non-null    bool
55 archstyle_craftsman                 4289 non-null    bool
56 archstyle_other                     4289 non-null    bool
57 archstyle_colonial                  4289 non-null    bool
58 archstyle_spanish                   4289 non-null    bool
59 archstyle_manufacturedhome-singlewide 4289 non-null    bool
60 archstyle_a-frame                   4289 non-null    bool
61 lotsizearea_listed_negative         4289 non-null    bool
62 lotsizearea_small                   4289 non-null    bool
63 originallistprice_persqft            4289 non-null    float64
64 originallistprice_scaled             4289 non-null    float64
65 originallistprice_persqft_scaled     4289 non-null    float64
dtypes: bool(43), datetime64[ns](1), float64(13), int64(6), object(3)
memory usage: 1.1+ MB

```

## Splitting the Data

Before examining interactions between various features and the target, I split the data into train and test sets, so as to avoid data leakage.

```
In [4]: train, test = split_data(df)
```

## Exploration

In this section, I look at individual features and various combinations of features and examine how they relate to the target variable of `newconstructionyn` in an attempt to identify which might be useful in identifying new construction listings.


*Note:* For statistical rigor, hypothesis testing should be conducted on many of the

exploratory findings shown here. For reasons of expediency, I did not perform statistical hypothesis tests as part of this analysis.

## Build Year

The most obvious feature to examine when attempting to identify new construction in the year the structure was built. If that year is recent, it seems fairly obvious that the property is more likely to be "new". There may be rare instances where a property has remained vacant for multiple years since it's construction, and there may be instances where a newly built home has it's build year input incorrectly in the listings, but these situations should be relatively rare.

```
In [5]: explore.build_year(train, target)
```



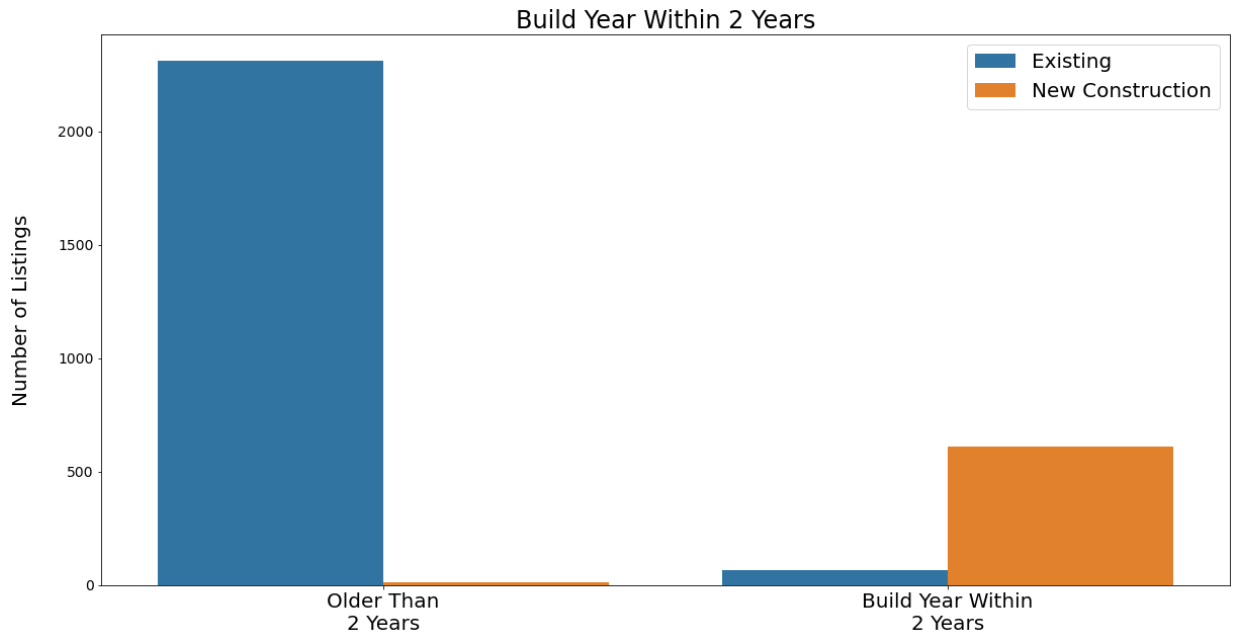
Here we can see that a build year of 2021 or 2022 is far more likely to be new construction,

and that new construction is more likely to have a build year of 2021 or 2022. This makes sense, because our dataset consists of listings from 2021 and 2022.

I also compared build year relative to the listing date, and created a binary column to determine if the build year was within two years of the listing date.

In [6]:

```
explore.built_last_two_years(train, target)
```



This leads me to believe that whether a property was built within two years of the listing date is likely to be the simplest, if not the most effective, way of determining whether the property is new construction.

## List Price

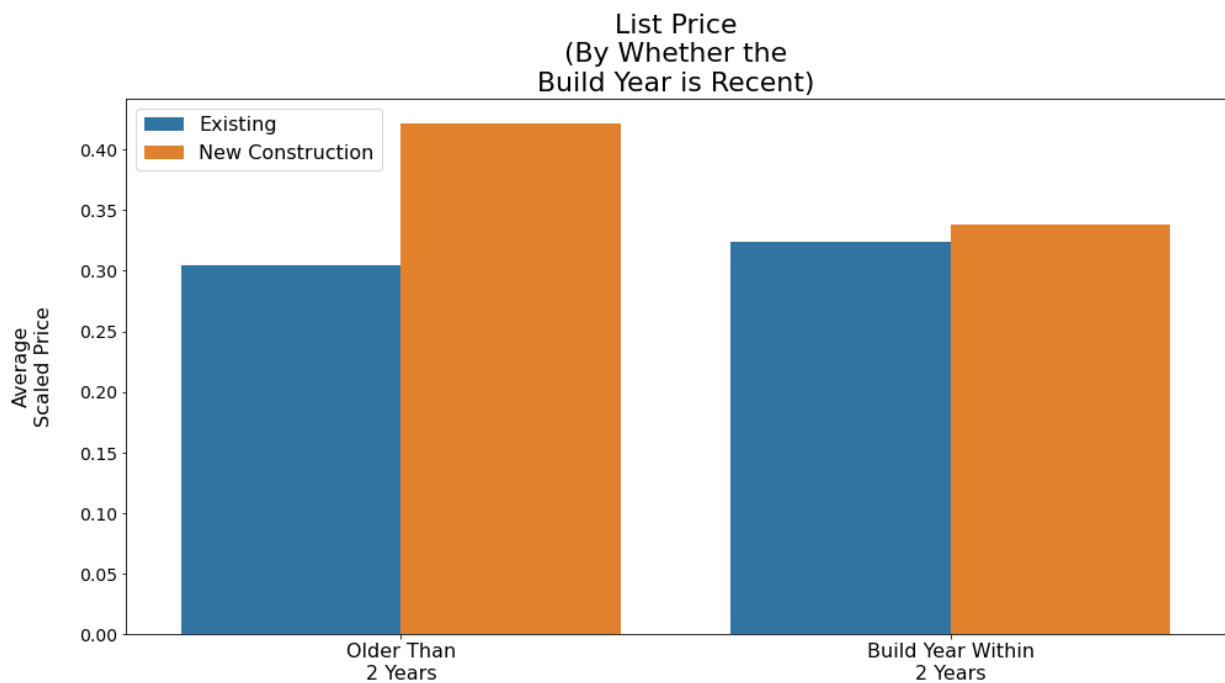
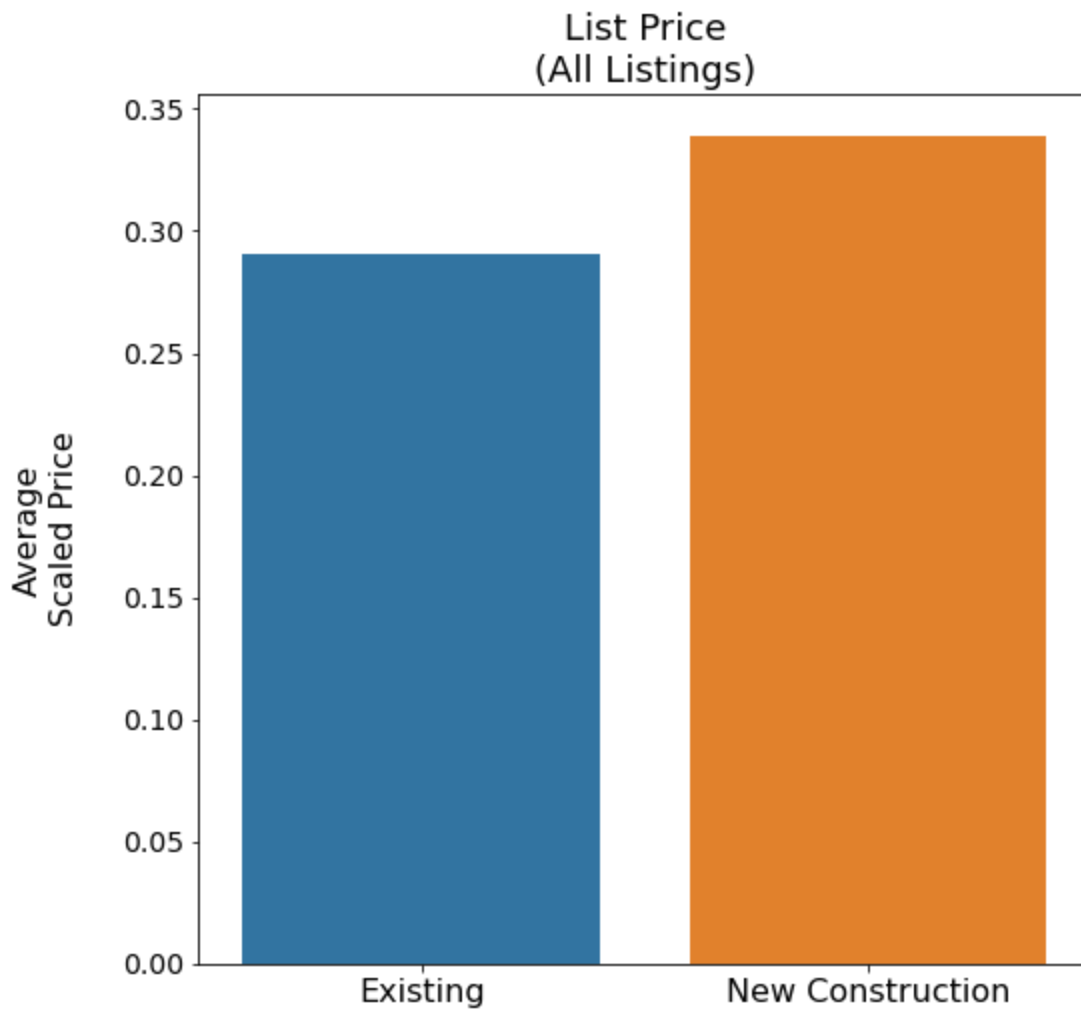
I'll hypothesize that new construction is generally more desirable, and will therefore have a higher price than a comparable home.

*Note:* Since the dataset includes both monthly rent prices and for-sale prices, these categories have each been scaled based on their respective min and max values.

In [7]:

```
explore.list_price(train, target)
```





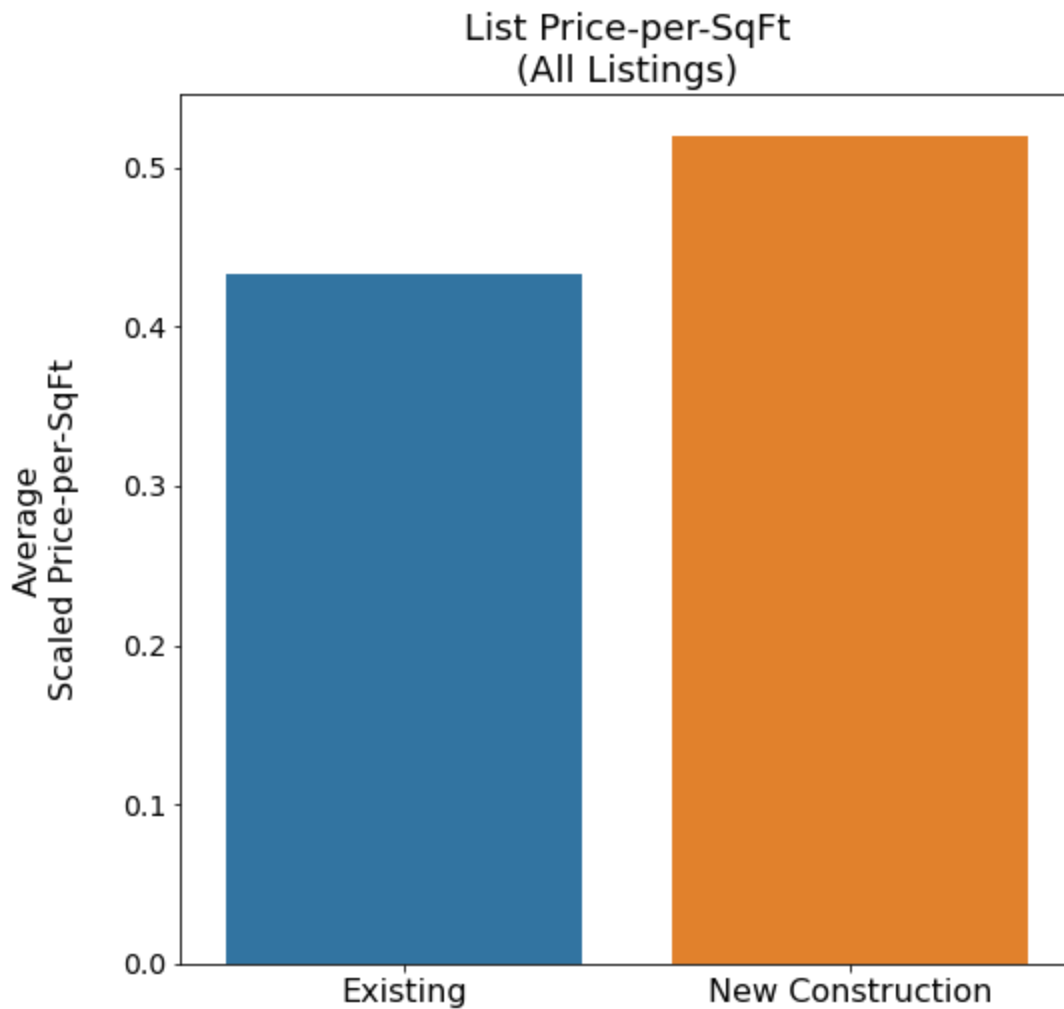
It does appear that new construction has a higher average list price, and that this holds true when controlling for whether the build year is recent. Though, the effect is lessened for recent build years, and it should be noted that there is a rather small sample size of new construction listings in older build years.

(With a larger dataset (and more time) we could group similar houses by other features in addition to build year, and see if a newly constructed home still holds more value than an existing home within those groups.)

## List Price - Per SqFt

Here we make a similar comparison using the per-square-foot price

```
In [8]: explore.list_price_persqft(train, target)
```





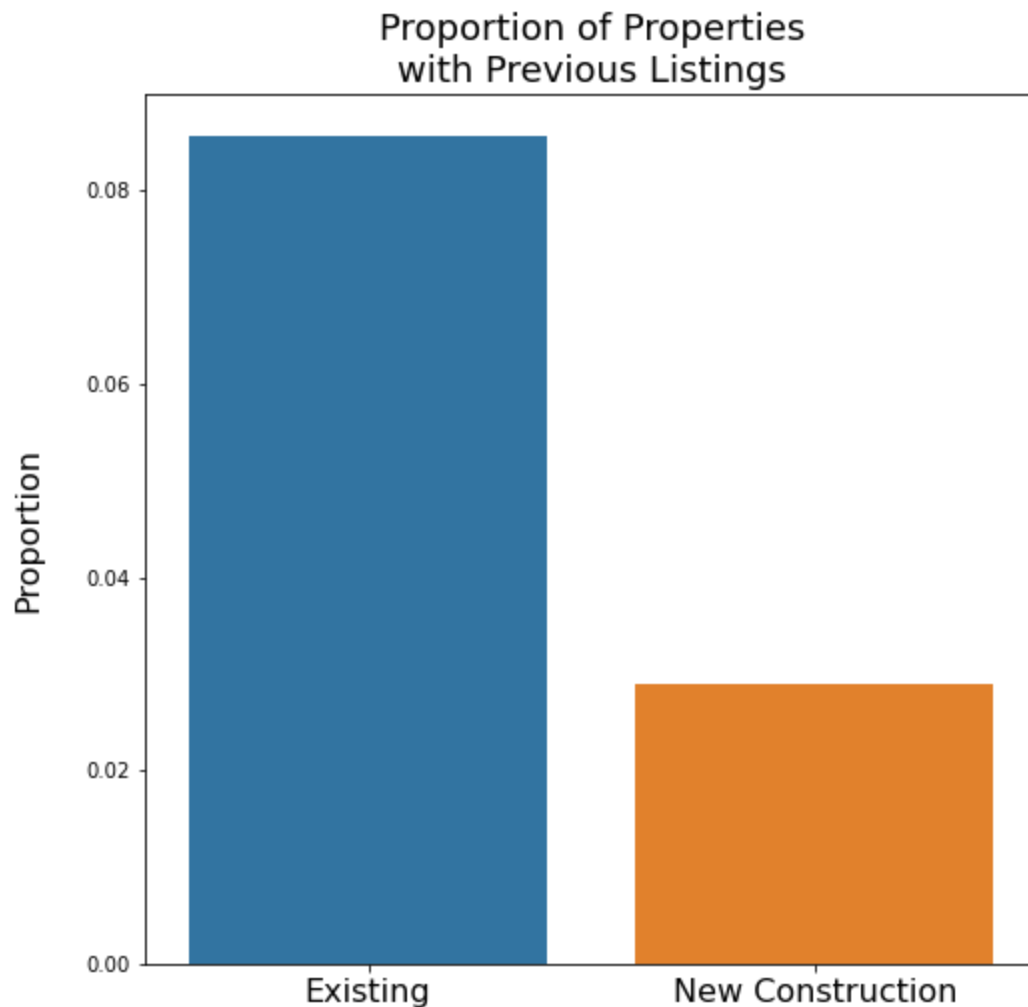
We see a very similar effect when comparing both the List Price and the List Price-per-SqFt between new construction vs existing homes. Since they are bound to be highly correlated, only one of these features should be included in any eventual model. I will choose to use the per-sqft price, since this is traditionally a better metric for comparing properties of differing size.

## Previous Listings

It was observed that many of the properties (as identified by their `address_id`) had multiple listings in the dataset. One likely explanation from this is that the property was removed from the MLS and then placed back on the MLS shortly thereafter, without a transaction taking place. However, one other explanation is that the earlier listings for a property represent a previous transaction. In this case, by definition, the more recent transactions would not be "new construction".

Therefore, I created a feature representing whether a property had a previous listing in the dataset, and hypothesize that new construction is less likely to have had a previous listing.

```
In [9]: explore.previous_listings(train, target)
```



It does the appear that new construction is much less likely to have had a previous listing in the dataset, though from this analysis we cannot know whether that is because the property did in fact have a previous transaction, or for some other reason. A more reliable way of using such information would be to compare the listing to a larger set of data going back further historically. For example, if an address had a previous listing from 5-10 years ago, it is probably even less likely to be new construction. We are limited by this dataset from this kind of analysis, since the dataset spans only about 1 yaer.

## Other Factors

Some other hypothesis were explored and found likely to be insignificant, so I will omit detailed visualizations from this report. Further information about these analyses can be found in the `explore_2.ipynb` in the drafts folder of this project's repository.

Some of the factors I considered were:

### Lot Size

Does New Construction have a different average lot size than existing homes?

While it did appear at first glance that new construction had a smaller lot size on average,

the difference appeared relatively weak.

### Central Heating and Cooling

One might assume that a new build would be much more likely than an older home to have a central HVAC system. However, there were so few observations at all which lacked central heating or cooling, that this did not appear to be a significant feature.

### Window/Wall Units

Similar to the central HVAC hypothesis discussed above, it seems new builds would be unlikely to have window/wall units. But very few homes at all in this dataset had them.

### Garages

One might notice driving around town that homes in some neighborhoods are more likely to have carports or other parking features than garages, and hypothesize that this is due to changing historical trends, and therefore new builds might be more (or less) likely to have a garage. But again, there were so few properties in the dataset without garages, that this did not appear to be a useful feature.

### Others

In a quick attempt to identify any other standout features or combinations of features, I used a seaborn pairplot, controlling for the target variable by changing it's color with the `hue` parameter. However, no major differences in new construction vs existing were visually apparent.

## Modeling

### Prep for Modeling

This function drops columns which will not be used in modeling, then encodes, splits, and scales data as appropriate.

```
In [10]: train, test = model.prep_for_modeling(df)
```

Now I will attempt to use the information learned during exploration to create a model capable of identifying at the time a property is listed whether it is new construction.

## Baseline Predictions

### Baseline 1

Since the majority of properties are not new construction, the simplest baseline would be to predict that each property is not new construction. This will maximize the accuracy of our baseline predictions.

## Baseline 2

Since I assume the goal is to avoid letting new construction go by unnoticed, I will attempt to optimize for the recall metric. In this case, a better baseline would be to treat all properties as if they are new construction.

## Model 1: A Simple 1-Feature Model

Exploration found that our most reliable feature is whether the build year of the property is within two calendar years of the listing date. This is also very intuitive - any reasonable person looking at a listing and trying to guess whether it was new construction would look for a recent build year. If our eventual model cannot predict more reliably than that, it will not be of much use. Therefore, Model 1 will make these simple predictions, and we will compare additional iterations of models against the performance of Model 1.

```
In [11]: # this function makes predictions for Model 1 and each of our baselines and e
model.baseline_models(train, target)
```

```
Out[11]:
```

	model_number	1	baseline_1	baseline_2
	metric_type			
	accuracy	0.974017	0.792472	0.207528
	f1_score	0.940092	0.000000	0.343724
	precision	0.901325	0.000000	0.207528
	recall	0.982343	0.000000	1.000000

## ML Models

Using the features which I identified as potentially useful during exploration, I created various types of machine learning classifiers.

In each case, I used sklearn's GridSearchCV to test varied combinations of hyperparameters and to cross-validate the results.

```
In [12]: features = ['built_last_two_years',
                    'previously_listed',
                    'scaled_stories',
                    'cooling_windowwall',
                    'originallistprice_persqft_scaled'
                    ]
```

## Decision Tree

```
In [13]: model.decision_tree(train, target, features)
```

```
Out[13]:
```

	max_depth	model_type	features	accuracy	recall	precision	F1_score
--	-----------	------------	----------	----------	--------	-----------	----------

	max_depth	model_type	features	accuracy	recall	precision	F1_score
0	2	decision tree	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.901575	0.940145
1	3	decision tree	[built_last_two_years, previously_listed, scal...	0.974682	0.974682	0.905698	0.941562
2	4	decision tree	[built_last_two_years, previously_listed, scal...	0.975348	0.975348	0.910573	0.942777
3	5	decision tree	[built_last_two_years, previously_listed, scal...	0.975015	0.975015	0.910410	0.941951
4	6	decision tree	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.908771	0.938782
			[built last two years				

Random Forest

In [14]:

model.random\_forest(train, target, features)

Out[14]:	max_depth	min_samples_leaf	model_type	features	accuracy	recall	prec
0	2	2	random forest	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.90
1	2	3	random forest	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.90
2	2	4	random forest	[built_last_two_years, previously_listed, scal...	0.971683	0.971683	0.90
3	3	2	random forest	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.90
4	3	3	random forest	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.90
5	3	4	random forest	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.90
6	4	2	random forest	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.90
7	4	3	random forest	[built_last_two_years, previously_listed, scal...	0.973350	0.973350	0.90
8	4	4	random forest	[built_last_two_years, previously_listed, scal...	0.973350	0.973350	0.90

	max_depth	min_samples_leaf	model_type	features	accuracy	recall	prec
9	5	2	random forest	[built_last_two_years, previously_listed, scal...	0.973350	0.973350	0.90
10	5	3	random forest	[built_last_two_years, previously_listed, scal...	0.973017	0.973017	0.90
11	5	4	random forest	[built_last_two_years, previously_listed, scal...	0.973017	0.973017	0.90
12	6	2	random forest	[built_last_two_years, previously_listed, scal...	0.973351	0.973351	0.90
13	6	3	random forest	[built_last_two_years, previously_listed, scal...	0.973350	0.973350	0.90
14	6	4	random forest	[built_last_two_years, previously_listed, scal...	0.973017	0.973017	0.90
15	7	2	random forest	[built_last_two_years, previously_listed, scal...	0.973684	0.973684	0.90
16	7	3	random forest	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.90
			random	[built_last_two_years,			

Logistic Regression

In [15]:

model.log\_regression(train, target, features)

	C	model_type	features	accuracy	recall	precision	F1_score
0	0.001	logistic regression	[built_last_two_years, previously_listed, scal...	0.792472	0.792472	0.000000	0.000000
1	0.010	logistic regression	[built_last_two_years, previously_listed, scal...	0.971354	0.971354	0.911647	0.932610
2	0.100	logistic regression	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.901575	0.940145
3	1.000	logistic regression	[built_last_two_years, previously_listed, scal...	0.974016	0.974016	0.901575	0.940145
4	10.000	logistic regression	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.901425	0.939331
5	100.000	logistic regression	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.901425	0.939331
6	1000.000	logistic regression	[built_last_two_years, previously_listed, scal...	0.973683	0.973683	0.901425	0.939331

Model Evaluation




Since none of our models improve upon Model 1's recall of 98.2%, our best method at present for predicting whether a property is new construction is simply to examine whether the build year of the listing is within two years of the listing date.


## Difference in Price?

To summarize the difference in price between new construction vs. existing, I will examine only the For-Sale listings, with the assumption that a valuation model is focused primarily on the purchase value of the property

### Total Price

```
In [16]: explore.difference_in_price(sale_df)
```





With the price being highly skewed, a median price is probably the best aggregate value.

```
In [17]: explore.show_median_prices(sale_df, target)
```

Median Price for New Construction: \$ 377,700


Median Price for Existing: \$ 305,000

Median difference in Price: \$ 72,700


The median difference in price between new Construction and Existing is around \$73,000 (for the city of San Antonio, June 2021 - July 2022)

## Price Per-SqFt

```
In [18]: explore.difference_in_price_persqft(sale_df)
```



### Distribution of Price-per-SqFt



```
In [19]: explore.show_median_price_persqft(sale_df, target)
```

Median Price-per-SqFt for New Construction: \$ 194

Median Price-per-SqFt for Existing: \$ 155

Median difference in Price-per-SqFt: \$ 39

The per-sqft prices are not as skewed, so an average price could also represent this data appropriately, but I will stick with median here for consistency.

The median difference in price-per-sqft between new construction and existing properties for this dataset is \$39.

```
In [ ]:
```